

## Coursebook: SQL Query using pandas

- Part 4 of Data Analytics Specialization
- Course Length: 9 hours
- Last Updated: July 2024 \_\_\_\_\_
- Author: Samuel Chan
- Developed by Algoritma's product division and instructors team

## Background

### Top-Down Approach

The coursebook is part 4 of the **Data Analytics Specialization** offered by Algoritma. It takes a more accessible approach compared to Algoritma's core educational products, by getting participants to overcome the “how” barrier first, rather than a detailed breakdown of the “why”.

This translates to an overall easier learning curve, one where the reader is prompted to write short snippets of code in frequent intervals, before being offered an explanation on the underlying theoretical frameworks. Instead of mastering the syntactic design of the Python programming language, then moving into data structures, and then the **pandas** library, and then the mathematical details in an imputation algorithm, and its code implementation; we would do the opposite: Implement the imputation, then a succinct explanation of why it works and applicational considerations (what to look out for, what are assumptions it made, when *not* to use it etc).

### Training Objectives

This coursebook is intended for participants who have completed the preceding courses offered in the **Data Analytics Developer Specialization**. This is the third course, **SQL and Data Visualization with Pandas**.

The coursebook focuses on:

- Querying from SQL Databases
- SQL Joins
- SQL Conditional Statements
- Flavors and Common Operators
- End to end data analysis

At the end of this course is a Graded Assignment section, where you are expected to apply all that you've learned on a new dataset, and attempt the given questions.

## Working with SQL Databases

There are a great number of python modules that provide functionalities to work with databases of all variants and flavors. For a MySQL database, we may form a connection using **pymysql** or one of many other alternatives:

```
import pymysql
conn = pymysql.connect(
    host=host,
    port=port,
    user=user,
    password=password,
    db=database)
```

We can then use `pd.read_sql_query()`, passing in the connection:

```
sales = pd.read_sql_query("SELECT * FROM sales", conn)
```

Under the hood, `pandas` uses `SQLAlchemy` so any database supported by that library will work. This isn't something you need to worry about at this stage of your learning journey, but for the sake of practice, let's also see how a connection URI for a `SQLite` database looks like:

```
import sqlite3
import pandas as pd

conn = sqlite3.connect("data_input/chinook.db")

albums = pd.read_sql_query("SELECT * FROM albums", conn)
albums.head()
```

##	AlbumId	Title	ArtistId
## 0	1	For Those About To Rock We Salute You	1
## 1	2	Balls to the Wall	2
## 2	3	Restless and Wild	2
## 3	4	Let There Be Rock	1
## 4	5	Big Ones	3

```
import sqlite3
import pandas as pd

conn = sqlite3.connect("data_input/chinook.db")

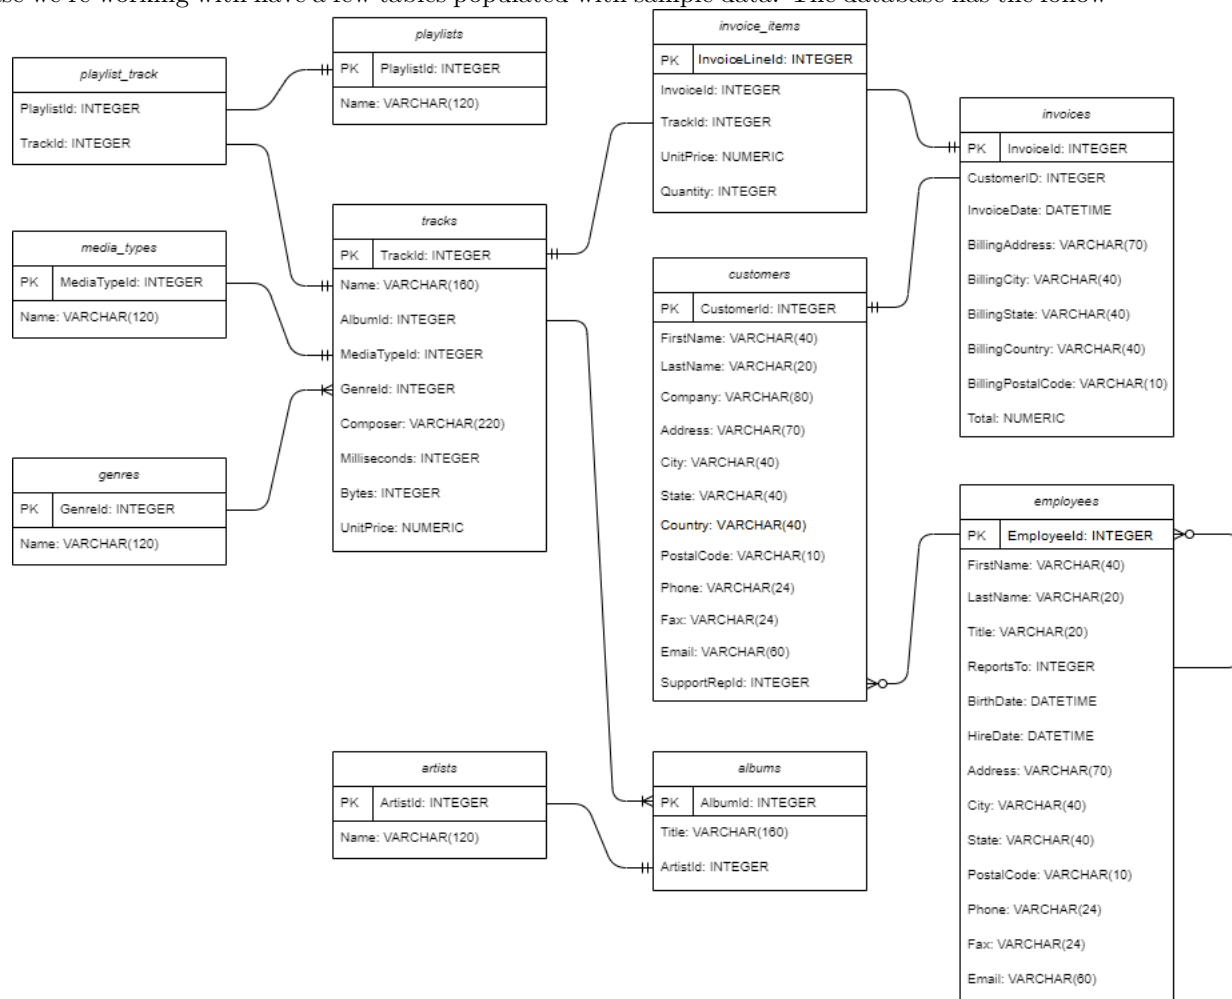
albums = pd.read_sql_query("SELECT * FROM employees", conn)
albums.head()
```

##	EmployeeId	LastName	...	Fax	Email
## 0	1	Adams	...	+1 (780) 428-3457	andrew@chinookcorp.com
## 1	2	Edwards	...	+1 (403) 262-3322	nancy@chinookcorp.com
## 2	3	Peacock	...	+1 (403) 262-6712	jane@chinookcorp.com
## 3	4	Park	...	+1 (403) 263-4289	margaret@chinookcorp.com
## 4	5	Johnson	...	1 (780) 836-9543	steve@chinookcorp.com
##					
##	[5 rows x 15 columns]				

In the above command, we asked for all columns of a table to be returned to us through the `SELECT *` command. Well, columns of which table? That would be `tables`. Together they form an SQL query:

```
SELECT * FROM albums
```

The database we're working with have a few tables populated with sample data. The database has the follow-



ing schema:

**Knowledge Check** We'll create a **DataFrame**: this time select all columns from the **artists** table. Recall that when we use `pd.read_sql_query()` command we pass in the SQL query as a string, and add a connection as the second parameter. Save the output as a **DataFrame**.

Your **DataFrame** should be constructed like this:

```
-- = pd.read_sql_query("SELECT * FROM artists", conn)
```

Question: 1. How many rows are there in your **DataFrame**?

```
## Your code below
```

```
## -- Solution code
```

The `pd.read_sql_query` is most commonly used with that two parameters above, but on its official documentation is a list of other parameters we can use as well.

In the following cell, we use a similar SQL query with an additional **LIMIT** statement to limit the output to the first 5 records (rows). However, notice that we also set `index_col` so the specified column is recognized as the index:


```
pd.read_sql_query("SELECT * FROM artists LIMIT 5",
                  conn,
                  index_col='ArtistId')
```

```
##                Name
## ArtistId
## 1                AC/DC
## 2                Accept
## 3                Aerosmith
## 4        Alanis Morissette
## 5        Alice In Chains
```

## SQL Joins

JOIN statements are used to combine records from two tables. We can have as many JOIN operations as we want in a SQL query.

Below is a diagram of the different types of SQL JOIN operations:



Credit: Data & Object Factory, LLC

In most business scenarios though, a **LEFT JOIN** is almost always the type of **JOIN** you want - it is very direct (and therefore easy to reason about). Left join return all records in the left table regardless if any of those records have a match in the right table.

The **INNER JOIN** is also very intuitive and easily understood. This query return all of the records in the left table that has a matching record in the right table.

As a personal side note, I've worked at companies where `RIGHT JOIN` is outright forbidden in favor of `LEFT JOIN`: directness and ease-of-understanding aside, all right joins can be replaced by the opposite left join.

The `OUTER JOIN` (also referred to as `FULL JOIN`) is also quite uncommon in practice. Performance reason aside, an outer join return all of the records from both tables regardless if there is a match or not, resulting in a DataFrame that has potentially a lot of `NULL` values.

Consider the database schema illustration again and pay attention to two tables and their respective columns:

1. `albums`:
  - `AlbumId, Title, ArtistId`
2. `artists`: `-ArtistId, Name`

We want a pandas DataFrame containing the `AlbumId`, `Title` and `Name`. Notice that `Name` is from the `artists` table while the other columns are from the `albums` table. What is a reasonable strategy?

The most straightforward solution is the `LEFT JOIN`, let's see an example:

```
albums = pd.read_sql_query("SELECT AlbumId, Title, a.Name \
                           FROM albums \
                           LEFT JOIN artists as a \
                           ON a.ArtistId = albums.ArtistId", conn)
albums.head()
```

	AlbumId	Title	Name
## 0	1	For Those About To Rock We Salute You	AC/DC
## 1	2	Balls to the Wall	Accept
## 2	3	Restless and Wild	Accept
## 3	4	Let There Be Rock	AC/DC
## 4	5	Big Ones	Aerosmith

Notice that in the code above, we place a backslash (`\`) character so we have line continuation and the newline will be ignored. This allows SQL to treat the entire query string as if they were essentially one line.

```
pd.read_sql_query("SELECT * FROM albums", conn).head()
```

	AlbumId	Title	ArtistId
## 0	1	For Those About To Rock We Salute You	1
## 1	2	Balls to the Wall	2
## 2	3	Restless and Wild	2
## 3	4	Let There Be Rock	1
## 4	5	Big Ones	3

**Knowledge Check** Consider the database schema illustration again and pay attention to two tables and their respective columns:

1. `albums`: `AlbumId, Title, ArtistId`
2. `tracks`: `TrackId, Name, AlbumId, GenreId, ... UnitPrice`

### 3. genres: GenreId, Name

Create a `DataFrame` containing all columns from the `tracks` table; Additionally, it should also contain: - The `Title` column from the `albums` table - The `Name` column from the `artists` table - The `Name` column from the `genres` table

**Hint 1:** In your `SELECT` statement, you can use `SELECT tracks.* FROM TRACKS` to select all columns from the `TRACKS` table

**Hint 2:** When we write `SELECT tracks.Name as tracksName`, we are renaming the output column from `Name` to `tracksName` using a technique called column aliasing. You may optionally consider doing this for columns that share the same name across different tables

Set the `TrackId` column to be the index. The resulting `DataFrame` should have 11 columns.

Give your `DataFrame` a name: name it `tracks`. Perform EDA on `tracks` to answer the following question:

1. Use `tail()` to inspect the last 5 rows of data. Which genre is present in the last 5 rows of our `tracks` `DataFrame` (Check all that apply)?
  - ☐ Latin
  - ☐ Classical
  - ☐ Soundtrack
  - ☐ Pop
2. Apply `pd.crosstab(..., columns='count')`, `.value_counts()`, or any other techniques you've learned to compute the frequency table of Genres in your `DataFrame`. Which is among the top 3 most represented genres in the `tracks` `DataFrame`?
  - ☐ Latin
  - ☐ Classical
  - ☐ Soundtrack
  - ☐ Pop
3. Use `groupby()` on Artist Name and compute the `mean()` on the `UnitPrice` of each tracks. You will realize that most artists price their tracks at 0.99 (`mean`) but there are several artists where the `mean()` is 1.99. Which of the Artist has a mean of 0.99 `UnitPrice`:
  - ☐ The Office
  - ☐ Aquaman
  - ☐ Pearl Jam
  - ☐ Lost

```
## Your code below
```

```
## -- Solution code
```

## SQL Aggregation

Since you have learned various aggregation tools in `pandas` such as `.crosstab()`, `.pivot_table()`, and `.groupby()`. It is also common practice to perform aggregation function using SQL. Consider the following query:

```
top_cust = pd.read_sql_query("SELECT CustomerId, SUM(Total) as TotalValue, \
                             COUNT(InvoiceId) as Purchases \
                             FROM INVOICES \
                             GROUP BY CustomerId \
                             ORDER BY TotalValue DESC \
                             LIMIT 5", con=conn, index_col='CustomerId')

top_cust
```

##	TotalValue	Purchases
## CustomerId		
## 6	49.62	7
## 26	47.62	7
## 57	46.62	7
## 45	45.62	7
## 46	45.62	7

Notice how the query fetched the top 5 customers from all time from the `invoices` table grouped by unique `CustomerId`. We performed two aggregation functions: `SUM()` and `COUNT()`, each aggregating different column from the `invoices` table. At the end, the `ORDER BY` statement is added to order the table based on the `TotalValue` column. Do note that the aggregated columns needs to be a numeric as the available aggregated functions are: `SUM`, `AVG`, `COUNT`, `MIN`, and `MAX`.

**Knowledge Check** Edit the following code to find out the most popular `genres` from all invoice sales. Use different column to acquire the following information: Summation of `Total` sales, and number of tracks bought from the `Quantity` columns.

hint: The `Total` can be obtained from `UnitPrice` and `Quantity` column in `invoice_items`

```
top_genre = pd.read_sql_query("SELECT genres.GenreId, genres.Name, \
                             __ (invoice_items.quantity * invoice_items.unitprice) AS Total, \
                             SUM(invoice_items.__) AS ___ \
                             FROM tracks \
                             LEFT JOIN genres ON _____ \
                             LEFT JOIN invoice_items ON _____ \
                             GROUP BY _____ \
                             "
                             conn,
                             index_col='GenreId'
                             )
```

Question: 1. What are the top 5 genres that generated the most profit?

*## Your code below*

## WHERE statements

We've seen how to use do some of the most common SQL operations this far. In particular, we have:

- Learned how to write `SELECT` statements

- Use `index_col` in the `pd.read_sql_query()` method
- SQL Join operations
- Use SQL Aliases

In the following example, we'll look at one more technique in the SQL arsenal: the `WHERE` clause

A `WHERE` clause is followed by a **condition**. If we want to query for all invoices where country of the billing address is Germany, we can add a `WHERE` clause to our sql query string:

```
germany = pd.read_sql_query("SELECT * FROM invoices WHERE BillingCountry = 'Germany'", conn)
germany.head()
```

```
##      InvoiceId  CustomerId  ... BillingPostalCode  Total
## 0           1           2  ...           70174    1.98
## 1           6          37  ...           60316    0.99
## 2           7          38  ...           10779    1.98
## 3          12           2  ...           70174   13.86
## 4          29          36  ...           10789    1.98
##
## [5 rows x 9 columns]
```

`WHERE` conditions can be combined with `IS`, `AND`, `OR` and `NOT`. Supposed we want to create a `DataFrame` containing all invoices where the billing country is **not** Germany, we can do the following:

```
not_germany = pd.read_sql_query("SELECT * FROM invoices WHERE BillingCountry IS NOT 'Germany'", conn)
not_germany.head()
```

```
##      InvoiceId  CustomerId  ... BillingPostalCode  Total
## 0           2           4  ...           0171    3.96
## 1           3           8  ...           1000    5.94
## 2           4          14  ...           T6G 2C7    8.91
## 3           5          23  ...           2113   13.86
## 4           8          40  ...           75002    1.98
##
## [5 rows x 9 columns]
```

Do note that the `IS` operator is the same as its mathematical notation counterpart: `=`. Similar with most programming language, it also supports other mathematical operator such as `>`, `>=`, `<`, and `<=`.

Now let's try another approach by using `IN` operator that enables us to specify multiple values for comparison. For example we'd like to retrieve all invoices from **Canada** and **USA**:

```
america = pd.read_sql_query("SELECT * FROM invoices WHERE BillingCountry IN ('USA', 'Canada')", conn)
america.head()
```

```
##      InvoiceId  CustomerId  ... BillingPostalCode  Total
## 0           4          14  ...           T6G 2C7    8.91
## 1           5          23  ...           2113   13.86
## 2          13          16  ...           94043-1351  0.99
## 3          14          17  ...           98052-8300    1.98
## 4          15          19  ...           95014    1.98
##
## [5 rows x 9 columns]
```



**Knowledge Check** Edit the following code to include a `WHERE` clause. We want the returned `DataFrame` to contain only the `Pop` genre and only when the `UnitPrice` of the track is 0.99:

```
popmusic = pd.read_sql_query("SELECT tracks.*, genres.Name as GenreName \
                              FROM tracks \
                              LEFT JOIN genres ON _____ \
                              WHERE genres.Name = _____ AND _____,
                              conn,
                              index_col='TrackId'
                              )
```

Question: 1. How many rows are there in `popmusic`?

```
## Your code below
```

```
## -- Solution code
```

## Operating Dates

Continuing from the last `WHERE` statements, we can retrieve all invoices billed to the country `Germany`. However, it is also common to perform a conditional query statement to retrieve a specific data range. Let's take a look at our `germanydata` frame for example:

```
germany.dtypes
```

```
## InvoiceId          int64
## CustomerId        int64
## InvoiceDate        object
## BillingAddress     object
## BillingCity        object
## BillingState       object
## BillingCountry     object
## BillingPostalCode  object
## Total            float64
## dtype: object
```

Notice how our `InvoiceDate` is listed as `Object` types. The `pd.read_sql_query` behaves like `pd.read_csv` where, by default, it reads data as numeric and object. This doesn't necessarily mean the database is stored using string format (commonly known as `VARCHAR` in SQL databases). Take a look at the following table schema:

```
invoices_table = pd.read_sql_query("SELECT sql FROM sqlite_master \
                                    WHERE name = 'invoices'", conn)
print(invoices_table.loc[0,:].values[0])
```

```
## CREATE TABLE "invoices"
## (
##     [InvoiceId] INTEGER PRIMARY KEY AUTOINCREMENT NOT NULL,
##     [CustomerId] INTEGER NOT NULL,
```

```
##      [InvoiceDate] DATETIME NOT NULL,
##      [BillingAddress] NVARCHAR(70),
##      [BillingCity] NVARCHAR(40),
##      [BillingState] NVARCHAR(40),
##      [BillingCountry] NVARCHAR(40),
##      [BillingPostalCode] NVARCHAR(10),
##      [Total] NUMERIC(10,2) NOT NULL,
##      FOREIGN KEY ([CustomerId]) REFERENCES "customers" ([CustomerId])
##      ON DELETE NO ACTION ON UPDATE NO ACTION
## )
```

Note that according to different DBMS you are using, there are different ways to retrieve a table's schema. The query above is used to retrieve table schema from SQLite database. The DATETIME type is stored with the following format: YYYY-MM-DD HH:MI:SS.

It is often useful to understand the table schema of your database so you can perform the appropriate operation. Within our `invoices`' schema you can see some of useful information such as: - `InvoiceId` is listed as a primary key - `InvoiceDate` is stored as DATETIME - `CustomerId` is registered as foreign key to `customers` table

If you are not provided with the database's schema, take some time to study each table schema. Now consider the following case: We are reviewing Germany market of the last year sales and would like to retrieve all invoices from the year 2012.

```
germany_2012 = pd.read_sql_query("SELECT * FROM invoices \
                                WHERE InvoiceDate >= '2012-01-01' AND InvoiceDate <= '2012-12-31'",
                                con=conn, parse_dates='InvoiceDate')
germany_2012['InvoiceDate'].describe()
```

```
## count                83
## mean      2012-07-04 15:02:10.120482048
## min                2012-01-01 00:00:00
## 25%                2012-03-31 12:00:00
## 50%                2012-06-30 00:00:00
## 75%                2012-09-29 12:00:00
## max                2012-12-30 00:00:00
## Name: InvoiceDate, dtype: object
```

Now there are other common approach using the `BETWEEN` operator, try and copy-paste the following code to a code cell:

```
germany_2012 = pd.read_sql_query("SELECT * FROM invoices \
                                WHERE InvoiceDate BETWEEN \
                                '___' AND '___'",
                                con=conn, parse_dates='InvoiceDate')
germany_2012['InvoiceDate'].describe()
```

Try completing the code above and see if the query fetch the same result as the previous one:

```
# Your code below
```

## Using LIKE Operator

Using a WHERE statement in a query can be beneficial to pull relevant data for our analysis. A common operator for a WHERE statement is LIKE. Consider the following SQL Query:

```
germany = pd.read_sql_query("SELECT * FROM invoices \
                             WHERE BillingCountry = 'Germany' \
                             AND BillingPostalCode LIKE '107%'", conn)
germany.head()
```

```
##      InvoiceId  CustomerId  ... BillingPostalCode  Total
## 0           7           38  ...             10779    1.98
## 1          29           36  ...             10789    1.98
## 2          30           38  ...             10779    3.96
## 3          40           36  ...             10789   13.86
## 4          52           38  ...             10779    5.94
##
## [5 rows x 9 columns]
```

A LIKE operator is a can be used to match a certain part of the data and can be really useful if we need to perform a partial matching to a specific value rather than using an equal-to operator. The 107% you see in the query means to extract value in BillingPostalCode that starts with the number 107. This is really helpful when you are aiming to extract data specifically on specific region. In Germany, you would know that Wilmersdorf and Tempelhof in Berlin has postal code starting with 107.

### Discussion:

If you were to put % on the end, you are matching everything that starts with 107, if you put it on the start %107, means you are matching everything that ends in 107. What do you think will came up if you use % before and after the pattern match?

**Dive Deeper:** Let's take a look at another case, consider the following data frame:

```
customerinv = pd.read_sql_query("SELECT firstname, lastname, email, company, \
                                 invoiceid, invoicedate, billingcountry, total \
                                 FROM invoices \
                                 left join customers \
                                 on invoices.customerId = customers.customerId", conn)
customerinv.head()
```

```
##      FirstName LastName  ... BillingCountry  Total
## 0      Leonie   Köhler  ...      Germany    1.98
## 1      Bjørn   Hansen  ...      Norway     3.96
## 2       Daan  Peeters  ...      Belgium     5.94
## 3       Mark  Philips  ...      Canada     8.91
## 4       John   Gordon  ...       USA     13.86
##
## [5 rows x 8 columns]
```

Within the first 6 data, you could see the Company column may not be reliable since most of them are filled with None. But if you pay close attention to the Email column, you could see some people have an email domain at apple, that could be an indicator of their company.

1. How would your conditional **WHERE** statement would be like if you want to count the number of customers that are working at Apple Inc.?

Try to complete the following codes:

```
applecust = pd.read_sql_query("SELECT firstname, lastname, email, company, \
                               invoiceid, invoicedate, billingcountry, total \
                               FROM invoices \
                               left join customers \
                               on invoices.____ = customers.____ \
                               WHERE email like '____'", conn)

)
```

*## Your code below*

*## -- Solution code*

Based on the data queried, how many of the customers is working at Apple Inc.?

- ☐ 412
- ☐ 49
- ☐ 7
- ☐ 14

## HAVING Statement

We have learned about how to perform aggregation in SQL and the usage of the **WHERE** statement. However, it's important to note that you can't use the **WHERE** clause to filter rows based on conditions on an aggregated column.

In the following example, we'll explore another technique in SQL: the **HAVING** clause.

The key difference between **HAVING** and **WHERE** is that the **HAVING** clause is used after the aggregation function. If we want to retrieve customers with a total invoice amount from all time greater than 47, let's consider the query below:

```
top_cust_47 = pd.read_sql_query("SELECT CustomerId, SUM(Total) as TotalValue \
                                FROM INVOICES \
                                GROUP BY CustomerId \
                                HAVING TotalValue > 47 \
                                ORDER BY TotalValue DESC", con=conn, index_col='CustomerId')

top_cust_47
```

```
##           TotalValue
## CustomerId
## 6                49.62
## 26               47.62
```

Notice how the query retrieves customers with a total invoice amount greater than 47 from all time. We performed an aggregation function to find the TotalValue for each CustomerId. We used the **HAVING** clause to filter the rows based on certain conditions in the aggregated column.

**Knowledge Check** Edit the following code to include a **HAVING** clause. We want the dataframe to only contain composers who have sold more than 15 tracks.

```
top_5_composer = pd.read_sql_query("SELECT tracks.____, SUM(invoice_items.___) AS TrackSold\
FROM tracks \
JOIN ___ ON ____ = ____ \
GROUP BY tracks._____\
HAVING ____ \
ORDER BY ____ DESC ",
conn)
```

Question :

1. How many Composer having TrackSold more than 15 ?

```
## Your code below
```

```
## -- Solution code
```

## SQL Subquery

In some cases, you'd like to fill in some value for the query from the database itself as one of the condition. For example, recall how we retrieved all customers that has the most total invoice in the previous exercise. Say from the information, we'd like to retrieve all the top customers invoice. To do that we will construct a **WHERE** statement using **IN** operator utilizing a subquery to retrieve all top customers:

```
customerinv = pd.read_sql_query("SELECT invoices.* \
FROM invoices \
WHERE invoices.CustomerId IN ( \
SELECT c.CustomerId FROM Customers as c \
LEFT JOIN invoices as i on i.CustomerId = c.CustomerId \
GROUP BY c.CustomerId \
ORDER BY SUM(Total) DESC LIMIT 10)", conn)

customerinv.head()
```

```
##      InvoiceId  CustomerId  ... BillingPostalCode  Total
## 0           46           6  ...           14300  8.91
## 1          175           6  ...           14300  1.98
## 2          198           6  ...           14300  3.96
## 3          220           6  ...           14300  5.94
## 4          272           6  ...           14300  0.99
##
## [5 rows x 9 columns]
```

Notice the subquery is defined as follow:

```
SELECT c.CustomerId FROM Customers as c
LEFT JOIN invoices as i on i.CustomerId = c.CustomerId
GROUP BY c.CustomerId
ORDER BY SUM(Total) DESC LIMIT 10
```

If used on its own, the query will retrieve the top 10 customer's IDs based on the total summation of their purchases. This when then used as the condition using IN statement on the previous query. The IN, however, can also be used using a hard-coded value like the following:

```
customerinv = pd.read_sql_query("SELECT * \
                                FROM invoices \
                                WHERE InvoiceId IN (46, 175, 198)", conn)
customerinv.head()
```

```
##      InvoiceId  CustomerId  ... BillingPostalCode  Total
## 0           46           6  ...              14300   8.91
## 1           175           6  ...              14300   1.98
## 2           198           6  ...              14300   3.96
##
## [3 rows x 9 columns]
```

**Knowledge Check** Imagine you're being instructed to analyze invoices that consist of a fair amount of tracks within one purchase. Previously, you have known that within one invoice the users would see a total of 1 to 10 tracks per purchase. Using subquery techniques you have learned, create an SQL query that fetched all invoice along with its total tracks quantity of bigger than 10. Complete the following query:

```
SELECT *,
(SELECT ___ FROM invoice_items GROUP BY ___) as Quantity
FROM invoices
WHERE Quantity > 10
```

```
## Your code below
```

```
## -- Solution code
```

## SQL Execution Order

Up to this point, we have discussed various SQL queries, starting from **SELECT**, the simplest query for selecting which data to present, to **HAVING**, for subsetting grouped data. Although **SELECT**, for example, is frequently written at the beginning of the query, it is not executed first. Therefore, in this chapter, we will mainly discuss the SQL execution order to gain a deeper understanding of how a set of queries work together.

First, it is necessary to define from which source the data will be queried. Therefore, **FROM** and **JOIN** statements are executed first. This includes the subqueries associated with these statements. Next, if any condition is defined, then the **WHERE** statement will be executed afterward. The rows that satisfy this condition will be included, and those that do not will be discarded. Following this, when we need to group the data, the **GROUP BY** statement is executed. After grouping, if we want to filter the grouped data, the **HAVING** statement is used.

Finally, the **SELECT** statement is executed to determine which parts of the queried data we want to display. The subsequent statements are **ORDER BY** and **LIMIT**, which order and limit the data display, respectively.

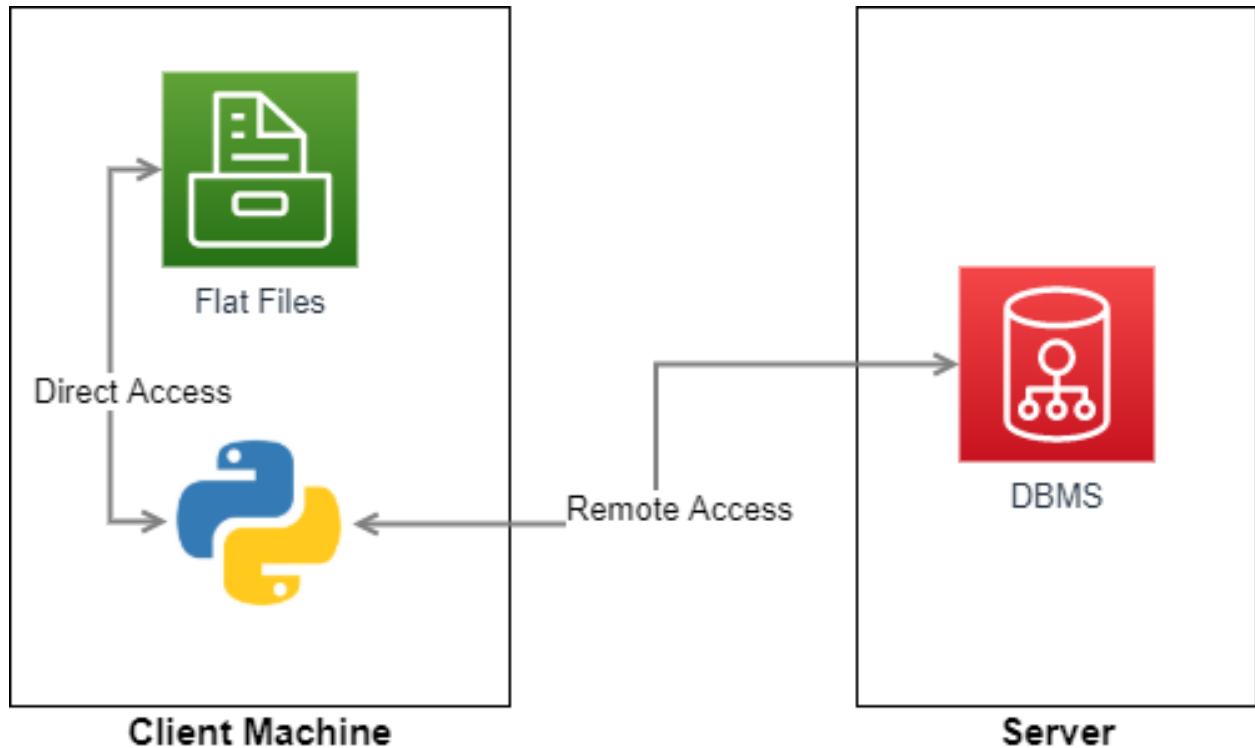
In summary, the SQL execution order is:

FROM - JOIN - WHERE - GROUP BY - HAVING - SELECT - ORDER BY - LIMIT.

## Under and Over Fetching

Now amidst all the tools selection available for data analysis, you will now ponder which one is better suited for you. To review let's recall what we have learned about Python features: - Reading flat files (csv / sas files) - Data cleansing and wrangling - Exploratory analysis tools - Visual exploratory tools

Now where does SQL fits in? First you will need to understand client-server architecture.



A bit different with Python when all your data operation is done on your local computer. When you worked with SQL, most likely you have a relational database stored remotely from your machine, usually a centralized database accessible to some clients.

When you perform a query, you execute a command to download the data to your computer. This downloading process require resources, and you need to effectively utilize the tools in order to minimize the overall cost.

### Discussion:

You are instructed to perform an analysis for all Rock genre sales on the last year (2012). Consider this questions:

- Is it necessary for you to download all tracks table to your computer?
- Will you filter the Rock genre tracks using SQL WHERE statements or conditional subsetting using Python?
- Since we need the information of multiple tables, which one is more convenient, querying a joined table or two separate tables from the database?

Try if you can construct your most optimum query below:

```
## Your code below
```

## Learn-by-Building

The following learn-by-building exercise will guide you through the process of building out a simple analysis along with some accompanying charts. This module is considerably more difficult than similar exercise blocks in the past, but it sure is a lot more rewarding!

Let's try by first constructing a DataFrame using the `read_sql_query()` method that we've grown familiar to. We want to develop a simple sales visualization report of our top 5 key markets (**Country** column in **customers**) ranked by Sales (**Total** column in **invoices**).

We also want to identify our top 5 customers by name (**FirstName**, **LastName**) in the report.

Last but not least, we want the report to include a day-of-week analysis on sales performance, and for that we will need the **InvoiceDate** column.

**Hint 1:** `pandas` has built-in methods of extracting the name of day in a week. We've seen this in Part 2 of this specialization (**Working with Datetime chapter**). An example usage is:

```
x['InvoiceDOW'] = x['InvoiceDate'].dt.weekday_name
```

**Hint 2:** In `read_sql_query`, you can use the `parse_dates='InvoiceDate'` argument to have the specified column parsed as date, saving you from a `to_datetime()` conversion

```
## Your code below
```

```
## -- Solution code
```